

```

1 import pandas as pd
2 from sklearn.preprocessing import LabelEncoder
3 from sklearn.model_selection import train_test_split
4 from sklearn.metrics import accuracy_score, classification_report, roc_curve, auc
5 import matplotlib.pyplot as plt
6 import seaborn as sns
7 import numpy as np
8 import tensorflow as tf
9 from tensorflow.keras.models import Sequential
10 from tensorflow.keras.layers import Dense, LSTM, Conv1D, MaxPooling1D, Flatten, Embedding

```

```

1 from tensorflow.keras.preprocessing.text import Tokenizer
2 from tensorflow.keras.preprocessing.sequence import pad_sequences
3 from tensorflow.keras.utils import to_categorical
4 from tensorflow.keras.callbacks import EarlyStopping

```

```

1 # Load the data from the file
2 file_path = 'breach_report_archive.txt'
3 data = pd.read_csv(file_path, delimiter='\t', encoding='latin1')

```

```

1 # Convert 'Breach Submission Date' to datetime
2 data['Breach Submission Date'] = pd.to_datetime(data['Breach Submission Date'])

```

```

1 # Encode categorical variables
2 label_encoders = {}
3 for column in ['State', 'Covered Entity Type', 'Type of Breach', 'Location of Breached Information', 'Business Associate Present']:
4     le = LabelEncoder()
5     data[column] = le.fit_transform(data[column])
6     label_encoders[column] = le

```

```

1 # Define categorical features and target variable
2 categorical_features = ['State', 'Covered Entity Type', 'Type of Breach', 'Location of Breached Information', 'Business Associate Present']
3 textual_feature = 'Web Description'
4 target = 'Individuals Affected'

```

```

1 # Binarize the target variable for classification
2 data['High Risk'] = data[target] > data[target].median()
3 y = data['High Risk'].astype(int)

```

```

1 # Tokenize and pad sequences for RNN and CNN models
2 max_words = 5000 # Maximum number of words to consider in the vocabulary
3 max_len = 512 # Maximum length of sequences

```

```

1 # Ensure the textual_feature column contains only strings
2 data[textual_feature] = data[textual_feature].astype(str)

```

```

1 tokenizer_keras = tf.keras.preprocessing.text.Tokenizer(num_words=max_words)
2 tokenizer_keras.fit_on_texts(data[textual_feature])
3 sequences = tokenizer_keras.texts_to_sequences(data[textual_feature])
4 X_textual_padded = tf.keras.preprocessing.sequence.pad_sequences(sequences, maxlen=max_len)

```

```

1 # Split the padded sequences into training and testing sets
2 X_train_textual, X_test_textual, y_train_textual, y_test_textual = train_test_split(X_textual_padded, y, test_size=0.2, random_state=42)

```

✓ RNN Model (LSTM)

```

1 from tensorflow.keras.layers import Dropout # Import Dropout

```

```

1 # Define RNN model (LSTM) with increased complexity and regularization
2 rnn_model = Sequential()
3 rnn_model.add(Embedding(input_dim=max_words, output_dim=128, input_length=max_len))
4 rnn_model.add(LSTM(128, return_sequences=True))
5 rnn_model.add(LSTM(64))
6 rnn_model.add(Dropout(0.5))

```

```
7 rnn_model.add(Dense(1, activation='sigmoid'))
8 rnn_model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
```

⚡ /usr/local/lib/python3.10/dist-packages/keras/src/layers/core/embedding.py:90: UserWarning: Argument `input_length` is deprecated. Just warnings.warn(
warnings.warn(
⏪ ⏩

```
1 # Train RNN model (LSTM) with early stopping
2 early_stopping = EarlyStopping(monitor='val_loss', patience=3, restore_best_weights=True)
3 rnn_model.fit(X_train_textual, y_train_textual, epochs=20, batch_size=32, validation_split=0.2, callbacks=[early_stopping])
```

⚡ Epoch 1/20
114/114 ————— 181s 1s/step - accuracy: 0.5365 - loss: 0.6882 - val_accuracy: 0.5895 - val_loss: 0.6729
Epoch 2/20
114/114 ————— 161s 1s/step - accuracy: 0.6310 - loss: 0.6506 - val_accuracy: 0.6970 - val_loss: 0.5947
Epoch 3/20
114/114 ————— 203s 1s/step - accuracy: 0.7805 - loss: 0.4866 - val_accuracy: 0.7387 - val_loss: 0.5430
Epoch 4/20
114/114 ————— 159s 1s/step - accuracy: 0.6587 - loss: 0.6060 - val_accuracy: 0.5609 - val_loss: 0.6858
Epoch 5/20
114/114 ————— 161s 1s/step - accuracy: 0.6585 - loss: 0.6108 - val_accuracy: 0.7530 - val_loss: 0.5349
Epoch 6/20
114/114 ————— 161s 1s/step - accuracy: 0.8179 - loss: 0.4194 - val_accuracy: 0.7596 - val_loss: 0.5320
Epoch 7/20
114/114 ————— 202s 1s/step - accuracy: 0.7595 - loss: 0.4787 - val_accuracy: 0.6158 - val_loss: 0.6121
Epoch 8/20
114/114 ————— 203s 1s/step - accuracy: 0.7015 - loss: 0.5539 - val_accuracy: 0.5335 - val_loss: 0.6893
Epoch 9/20
114/114 ————— 201s 1s/step - accuracy: 0.5083 - loss: 0.7002 - val_accuracy: 0.5357 - val_loss: 0.6896
<keras.src.callbacks.history.History at 0x7d56a4fbfd0>

```
1 # Evaluate RNN model (LSTM)
2 y_pred_rnn_proba = rnn_model.predict(X_test_textual).flatten()
3 y_pred_rnn = (y_pred_rnn_proba > 0.5).astype(int)
4 accuracy_rnn = accuracy_score(y_test_textual, y_pred_rnn)
5 report_rnn = classification_report(y_test_textual, y_pred_rnn)
```

⚡ 36/36 ————— 14s 377ms/step

✓ CNN Model

```
1 from tensorflow.keras.layers import GlobalMaxPooling1D # Import GlobalMaxPooling1D
```

```
1 # Define CNN model with increased complexity and regularization
2 cnn_model = Sequential()
3 cnn_model.add(Embedding(input_dim=max_words, output_dim=128, input_length=max_len))
4 cnn_model.add(Conv1D(128, 5, activation='relu'))
5 cnn_model.add(MaxPooling1D(pool_size=4))
6 cnn_model.add(Conv1D(64, 5, activation='relu'))
7 cnn_model.add(GlobalMaxPooling1D())
8 cnn_model.add(Dropout(0.5))
9 cnn_model.add(Dense(1, activation='sigmoid'))
10 cnn_model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
```

⚡ /usr/local/lib/python3.10/dist-packages/keras/src/layers/core/embedding.py:90: UserWarning: Argument `input_length` is deprecated. Just warnings.warn(
warnings.warn(
⏪ ⏩

```
1 # Train CNN model with early stopping
2 cnn_model.fit(X_train_textual, y_train_textual, epochs=20, batch_size=32, validation_split=0.2, callbacks=[early_stopping])
```

⚡ Epoch 1/20
114/114 ————— 38s 311ms/step - accuracy: 0.5494 - loss: 0.6810 - val_accuracy: 0.8277 - val_loss: 0.4044
Epoch 2/20
114/114 ————— 29s 255ms/step - accuracy: 0.8314 - loss: 0.3461 - val_accuracy: 0.8869 - val_loss: 0.2647
Epoch 3/20
114/114 ————— 41s 252ms/step - accuracy: 0.8969 - loss: 0.2179 - val_accuracy: 0.8858 - val_loss: 0.2695
Epoch 4/20
114/114 ————— 41s 257ms/step - accuracy: 0.9163 - loss: 0.1792 - val_accuracy: 0.8683 - val_loss: 0.3156
Epoch 5/20
114/114 ————— 41s 256ms/step - accuracy: 0.9261 - loss: 0.1398 - val_accuracy: 0.8705 - val_loss: 0.3186
<keras.src.callbacks.history.History at 0x7d569ae1c7c0>

```

1 # Evaluate CNN model
2 y_pred_cnn_proba = cnn_model.predict(X_test_textual).flatten()
3 y_pred_cnn = (y_pred_cnn_proba > 0.5).astype(int)
4 accuracy_cnn = accuracy_score(y_test_textual, y_pred_cnn)
5 report_cnn = classification_report(y_test_textual, y_pred_cnn)

```

↗ 36/36 ————— 2s 62ms/step

✓ Evaluate and Compare Models

Plot ROC Curves and Confusion Matrices

```

1 # Function to plot ROC curve
2 def plot_roc_curve(y_true, y_pred_proba, model_name):
3     fpr, tpr, _ = roc_curve(y_true, y_pred_proba)
4     roc_auc = auc(fpr, tpr)
5     plt.figure(figsize=(8, 6))
6     plt.plot(fpr, tpr, color='darkorange', lw=2, label=f'ROC curve (area = {roc_auc:.2f})')
7     plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
8     plt.xlim([0.0, 1.0])
9     plt.ylim([0.0, 1.05])
10    plt.xlabel('False Positive Rate')
11    plt.ylabel('True Positive Rate')
12    plt.title(f'Receiver Operating Characteristic (ROC) Curve for {model_name}')
13    plt.legend(loc="lower right")
14    plt.show()

```

```

1 # Import the necessary libraries
2 from sklearn.metrics import roc_curve, auc
3 from sklearn.metrics import confusion_matrix

```

```

1 # Function to plot confusion matrix
2 def plot_confusion_matrix(y_true, y_pred, model_name):
3     cm = confusion_matrix(y_true, y_pred)
4     plt.figure(figsize=(8, 6))
5     sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=['Low Risk', 'High Risk'], yticklabels=['Low Risk', 'High Risk'])
6     plt.xlabel('Predicted')
7     plt.ylabel('Actual')
8     plt.title(f'Confusion Matrix for {model_name}')
9     plt.show()

```


```

1 # Plot ROC curve and confusion matrix for RNN
2 plot_roc_curve(y_test_textual, y_pred_rnn_proba, 'RNN (LSTM)')
3 plot_confusion_matrix(y_test_textual, y_pred_rnn, 'RNN (LSTM)')

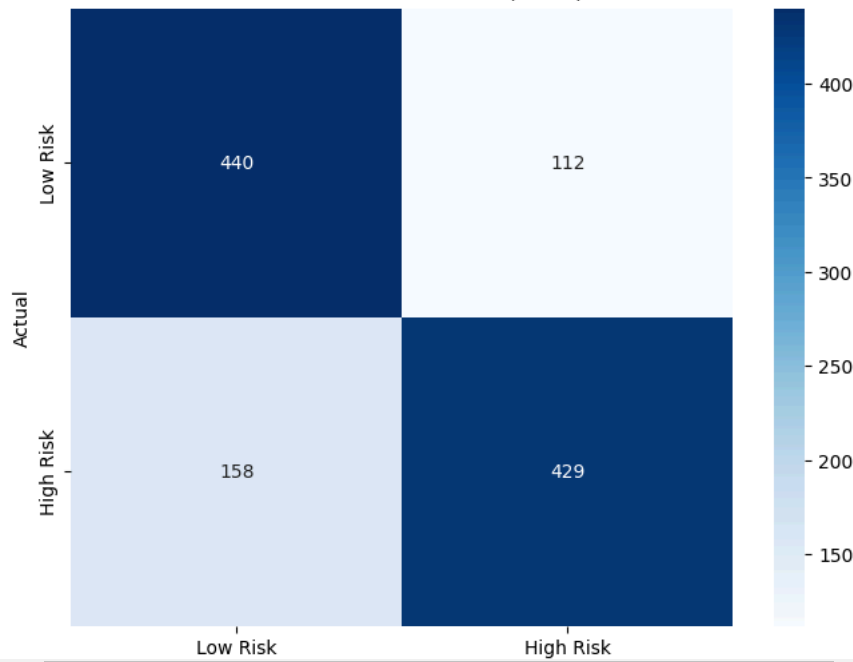
```



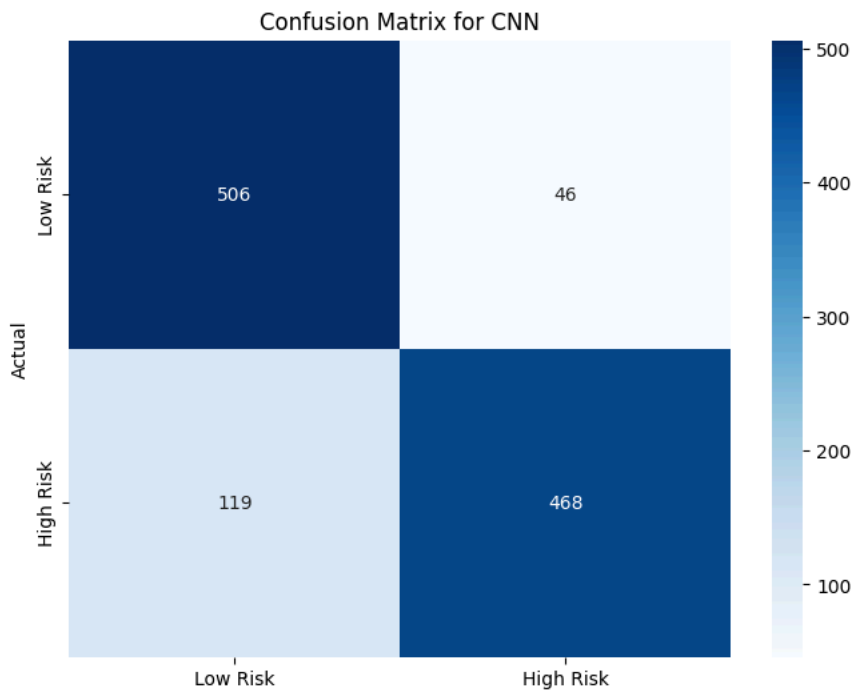
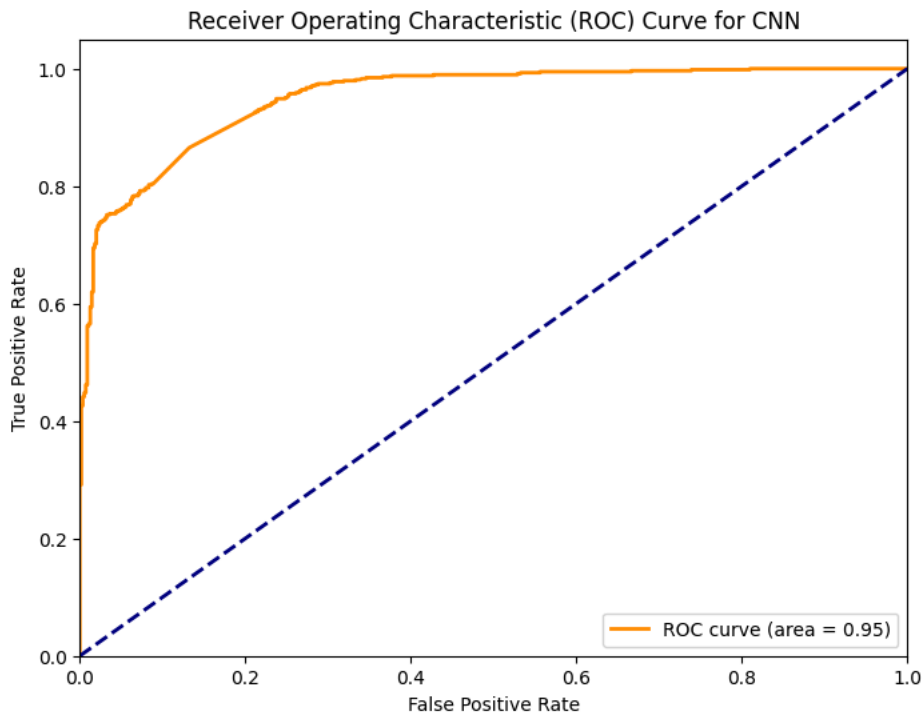
Receiver Operating Characteristic (ROC) Curve for RNN (LSTM)



Confusion Matrix for RNN (LSTM)



```
1 # Plot ROC curve and confusion matrix for CNN
2 plot_roc_curve(y_test_textual, y_pred_cnn_proba, 'CNN')
3 plot_confusion_matrix(y_test_textual, y_pred_cnn, 'CNN')
```



```
1 # Print results for RNN
2 print(f"Model: RNN (LSTM)")
3 print(f"Accuracy: {accuracy_rnn}")
4 print(f"Classification Report:\n{report_rnn}\n")
```



Model: RNN (LSTM)
Accuracy: 0.7629499561018437
Classification Report:

	precision	recall	f1-score	support
0	0.74	0.80	0.77	552
1	0.79	0.73	0.76	587
accuracy			0.76	1139
macro avg	0.76	0.76	0.76	1139
weighted avg	0.77	0.76	0.76	1139

```
1 # Print results for CNN
2 print(f"Model: CNN")
3 print(f"Accuracy: {accuracy_cnn}")
4 print(f"Classification Report:\n{report_cnn}\n")
```



Model: CNN

Accuracy: 0.8551360842844601

Classification Report:

	precision	recall	f1-score	support
0	0.81	0.92	0.86	552
1	0.91	0.80	0.85	587
accuracy			0.86	1139
macro avg	0.86	0.86	0.85	1139
weighted avg	0.86	0.86	0.85	1139